

# Operating Systems Principles

cosc1112/cosc1114

School of Science

Semester 2, 2017

## Lecture 10 – I/O System

Dr. Ke Deng

[ke.deng@rmit.edu.au](mailto:ke.deng@rmit.edu.au)

# Outline

- Overview
- I/O Hardware
- Application I/O Interface
- Kernel I/O Subsystem
- Transforming I/O Requests to Hardware Operations
- STREAMS
- Performance

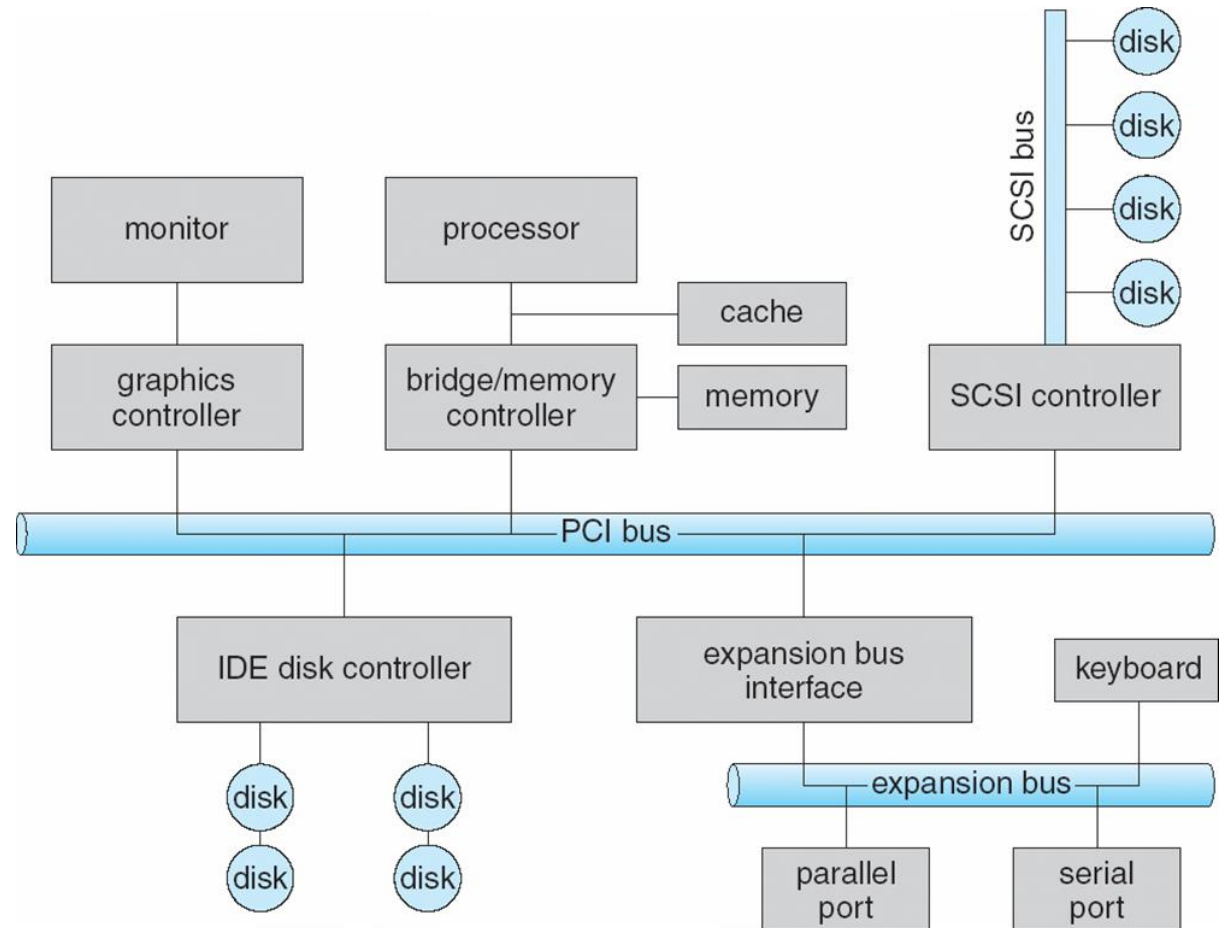
# Objectives

- Explore the structure of an operating system's I/O subsystem
- Discuss the principles of I/O hardware and its complexity
- Provide details of the performance aspects of I/O hardware and software

# Overview

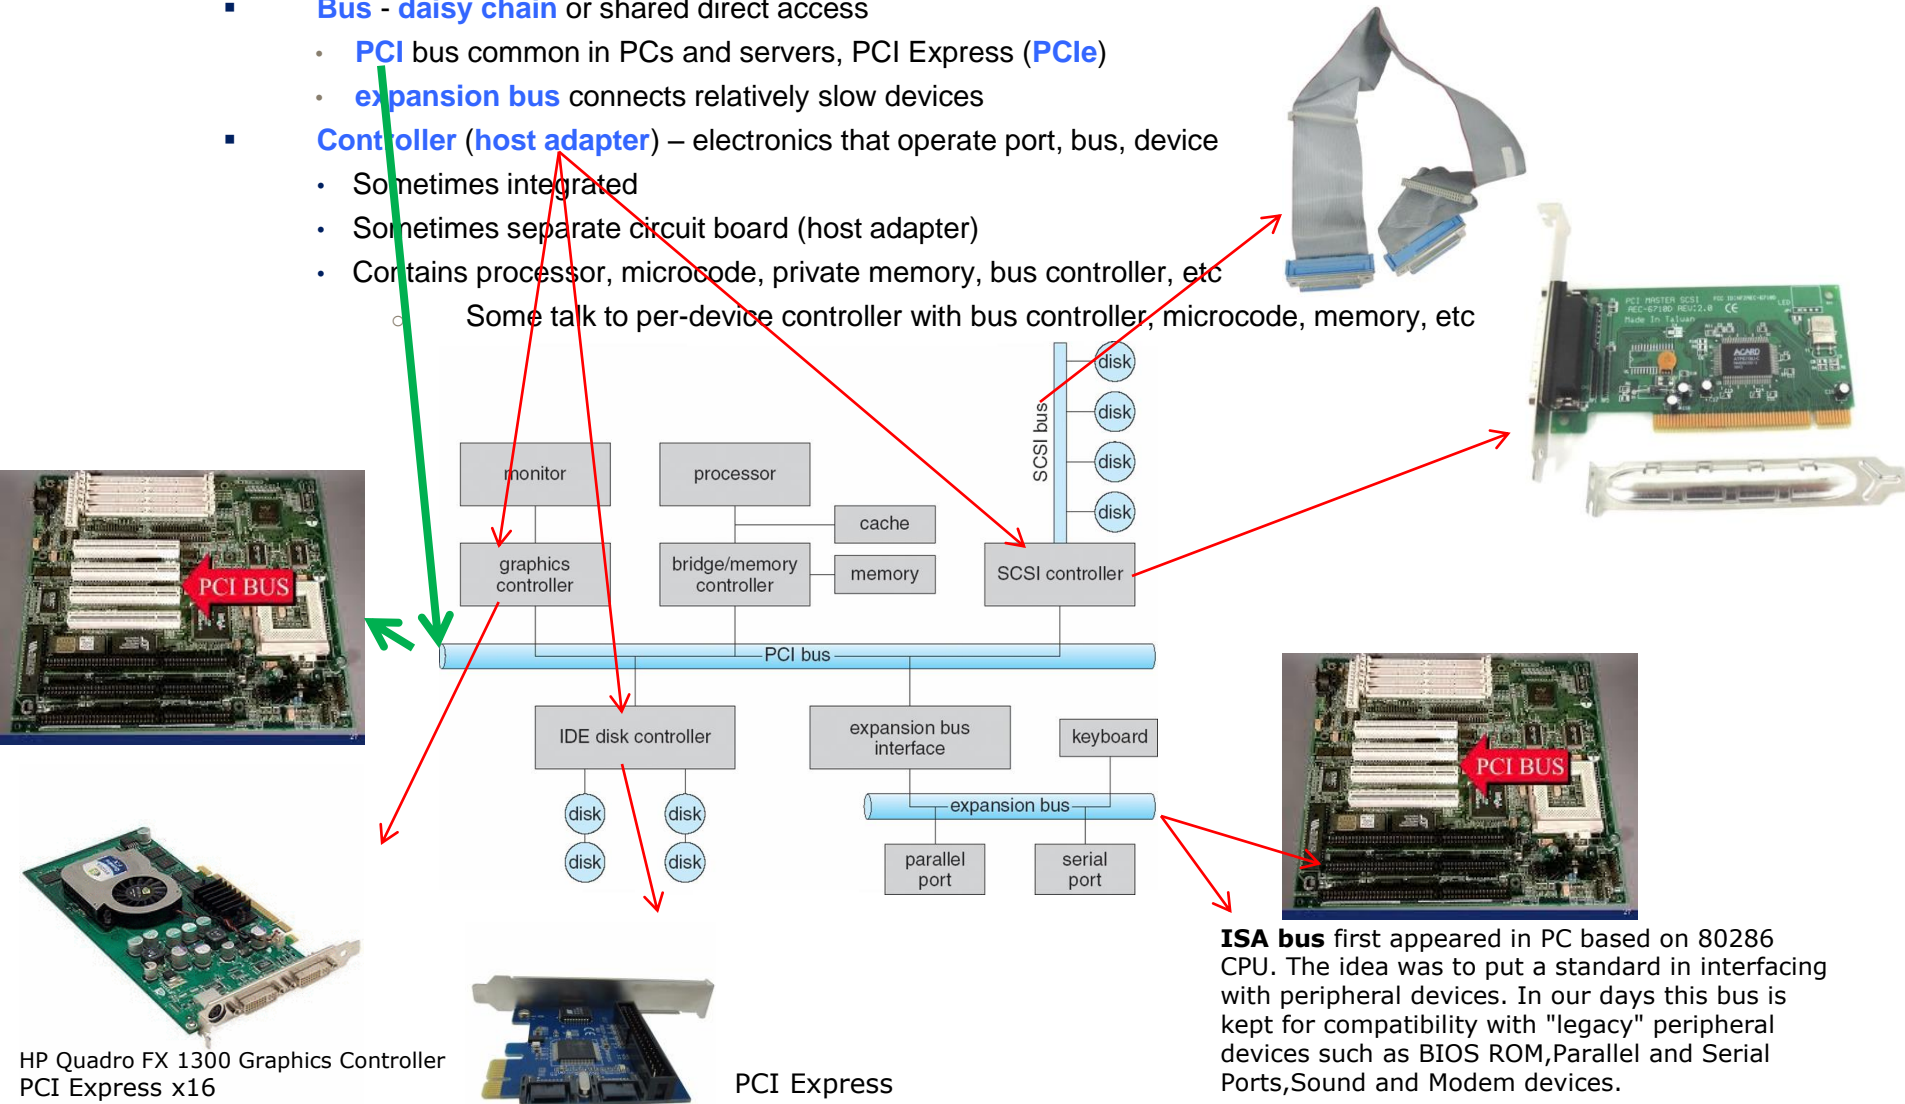
- I/O management is a major component of operating system design and operation
- Ports, busses, device controllers connect to various devices

**Device drivers**  
encapsulate  
device details -  
Present uniform  
device-access  
interface to I/O  
subsystem



## Common concepts – signals from I/O devices interface with computer

- **Port** – connection point for device
  - **Bus** - **daisy chain** or shared direct access
    - **PCI** bus common in PCs and servers, PCI Express (**PCIe**)
    - **expansion bus** connects relatively slow devices
  - **Controller (host adapter)** – electronics that operate port, bus, device
    - Sometimes integrated
    - Sometimes separate circuit board (host adapter)
    - Contains processor, microcode, private memory, bus controller, etc
- Some talk to per-device controller with bus controller, microcode, memory, etc



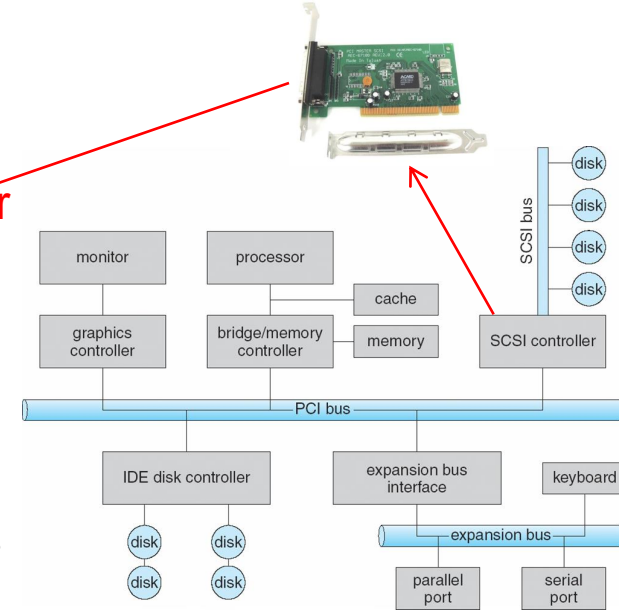


# I/O Hardware (Cont.)

How the processor give command and data to a controller to accomplish an I/O transfer?

- Devices have registers (typically 1-4 bytes in size, or FIFO buffer)
  - data-in register - read by the host to get input
  - data-out register - written by the host to send output
  - status register - read by the host to indicate such as current command has completed
  - control register – read by the host to start a command and change mode of a device (e.g. parity checking)

- Device driver communicate with registers by two ways
  - Direct I/O instructions
    - Use special I/O instructions that specify the transfer of a byte or word to an I/O port address. The I/O instruction triggers bus lines to select the proper device and to move bits into or out of a device register.
  - Memory-mapped I/O
    - Device control registers are mapped into the address space of the processor. The CPU executes I/O requests using the standard data-transfer instructions to read the write the device control-registers.
    - For example, the graphics controller has I/O ports for basic control operations; but the controller has a large-memory mapped region to hold screen contents. The process send outputs to the screen by writing data into the memory-mapped region.



I/O address range (hexadecimal)	device
000-00F	DMA controller
020-021	interrupt controller
040-043	timer
200-20F	game controller
2F8-2FF	serial port (secondary)
320-32F	hard-disk controller
378-37F	parallel port
3D0-3DF	graphics controller
3F0-3F7	diskette-drive controller
3F8-3FF	serial port (primary)

# Polling

## ■ For each byte of I/O

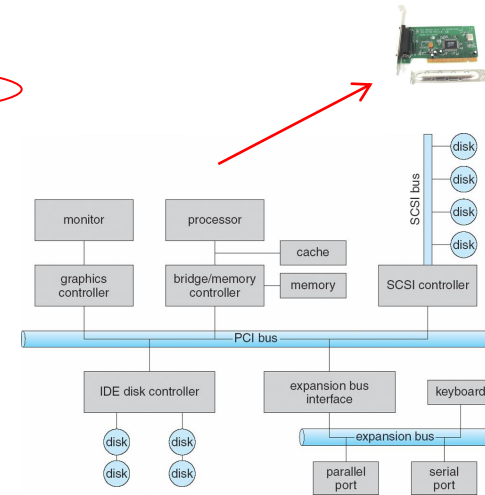
1. The host repeatedly reads the busy bit in status register until it becomes 0;
2. The host sets the write bit in the command register and write a byte into the data-out register;
3. The host sets command-ready bit in command register;
4. When the controller notices that the command-ready bit is set, it sets the busy bit;
5. The controller reads the command register and sees the write command, it reads the data-out register to get the bytes and does the I/O to the device;
6. when transfer done, device controller clears the busy bit, error bit in status register, command-ready bit in command register;

Step 1 is **busy-wait** or **polling** cycle to wait for I/O from device

Polling can happen in 3 instruction cycles

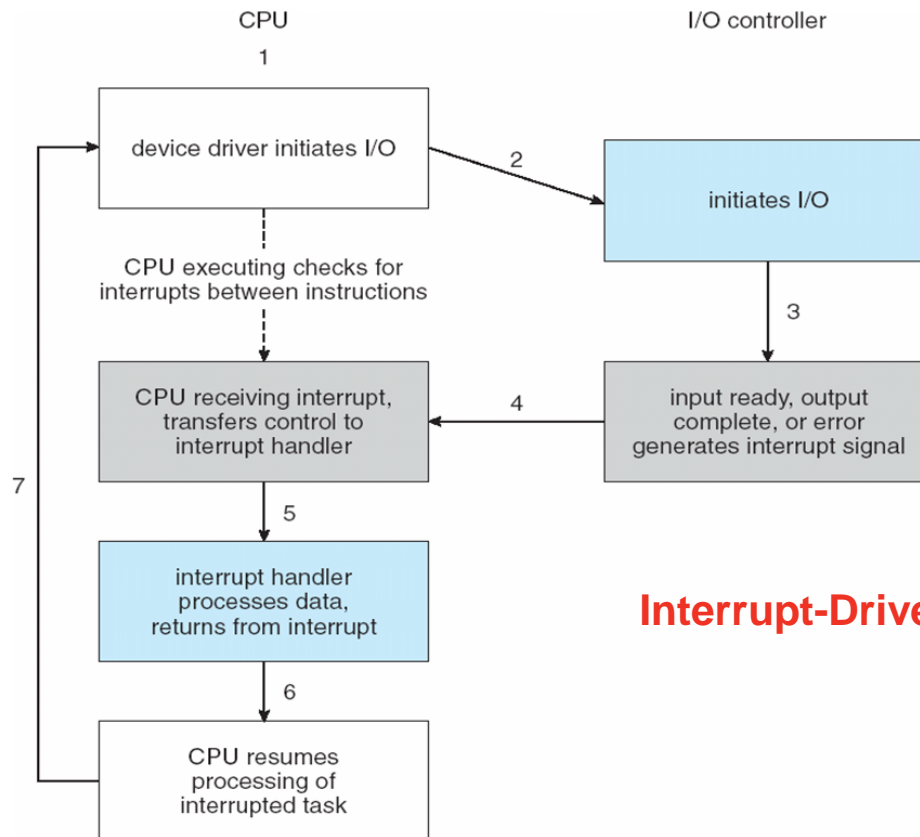
- Read status from status register
- logical-and to extract status bit
- branch if not zero

How to be more efficient if non-zero infrequently?



# Interrupts

- The CPU has a wire called **Interrupt-request line**
  - It is checked by processor after each instruction
  - When the CPU detects that a device controller has asserted a signal on the interrupted line, the CPU saves current state and jumps to the interrupt-handler routine at a fixed address in memory
  - The interrupt handler determined the cause of the interrupt, performing the necessary processing.
  - The interrupt handler returns the CPU to the execution state prior to the interrupt.



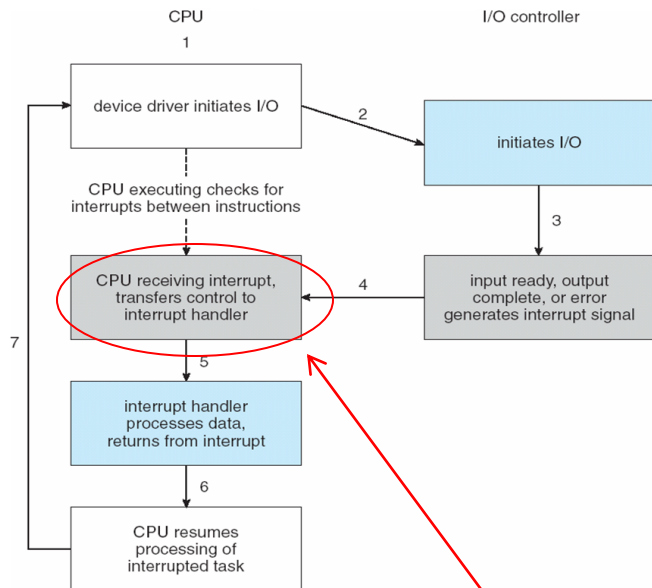
**Interrupt-Driven I/O Cycle**



# Interrupts (cont.)

## Intel Pentium Processor Event-Vector Table

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts



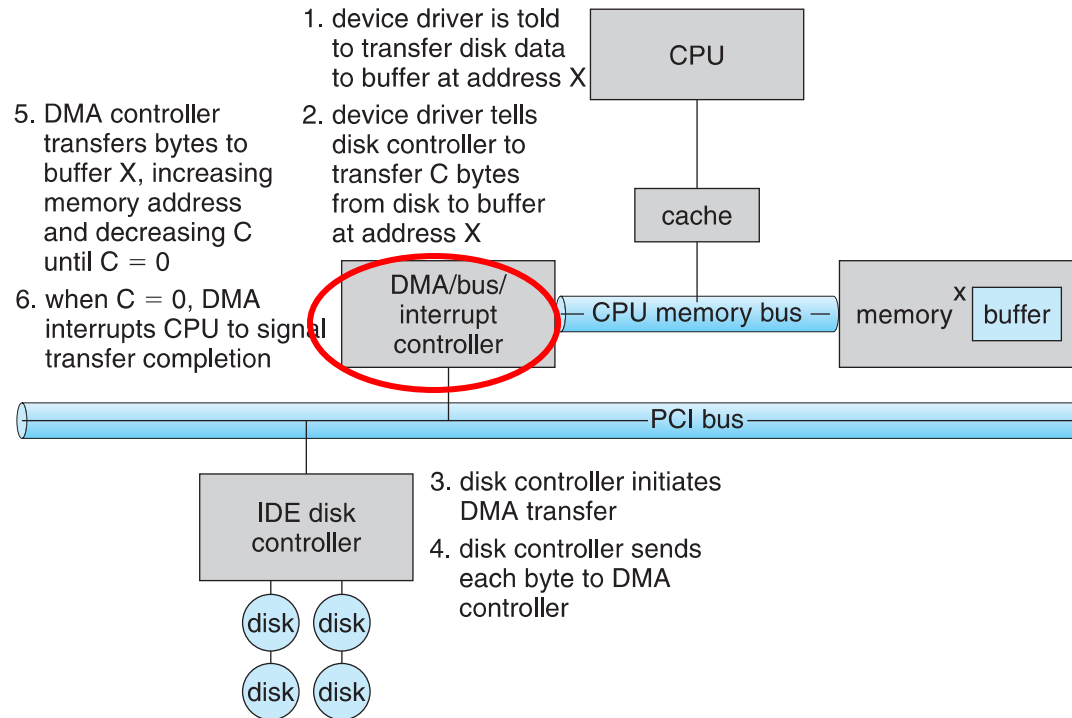
Interrupt vector contains the memory addresses of specialized interrupt handler.

# Interrupts (Cont.)

- Software generated interruption
  - Interrupt mechanism also used for **exceptions**  
Terminate process, crash system due to hardware error
  - **Page fault** executes when memory access error
  - System call executes via **trap** to trigger kernel to execute request
- Multi-CPU systems can process interrupts concurrently
  - If operating system designed to handle it

# Direct Memory Access (DMA)

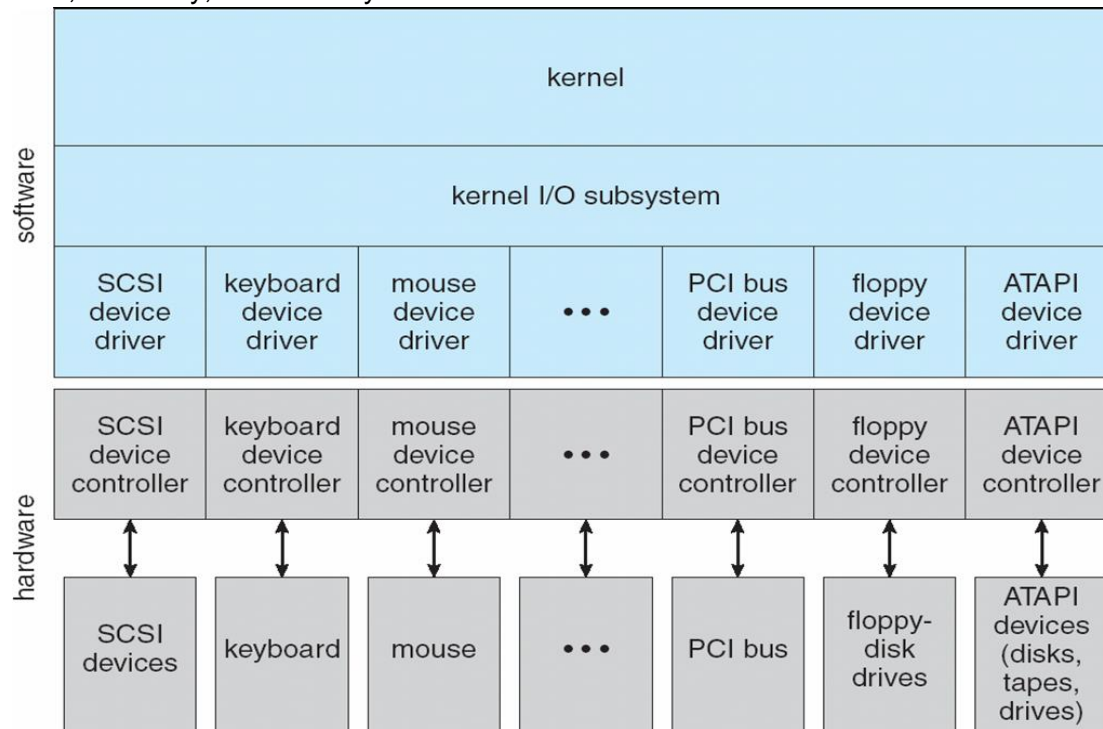
- For large data I/O, **DMA** controller can be used; Bypasses CPU to transfer data directly between I/O device and memory



- A simple DMA controller is a standard component in PCs.
- While DMA controller seizes the memory bus, the CPU is momentarily prevented from accessing main memory – [cycling stealing](#). But CPU can still access its primary and secondary caches.
- Some computer architectures perform [direct virtual memory access \(DVMA\)](#): translate the virtual memory address to physical memory.

# Application I/O Interface

- Device-driver layer hides differences among I/O controllers from kernel; Each OS has its own I/O subsystem structures and device driver frameworks
- Devices vary in many dimensions
  - Character-stream or block
  - Sequential or random-access
  - Synchronous or asynchronous (or both)
  - Sharable or dedicated
  - Speed of operation
  - read-write, read only, or write only



## Block and Character Devices

- Block devices include disk drives
  - Commands include read, write, seek
  - Raw I/O, direct I/O, or file-system access
  - Memory-mapped file access possible
    - File mapped to virtual memory and clusters brought via demand paging
  - DMA
- Character devices include keyboards, mice, serial ports
  - Commands include `get()`, `put()`
  - Libraries layered on top allow line editing

## Network Devices

- Varying enough from block and character to have own interface
- Linux, Unix, Windows and many others include **socket** interface
  - Separates network protocol from network operation
  - Includes `select()` functionality
- Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)

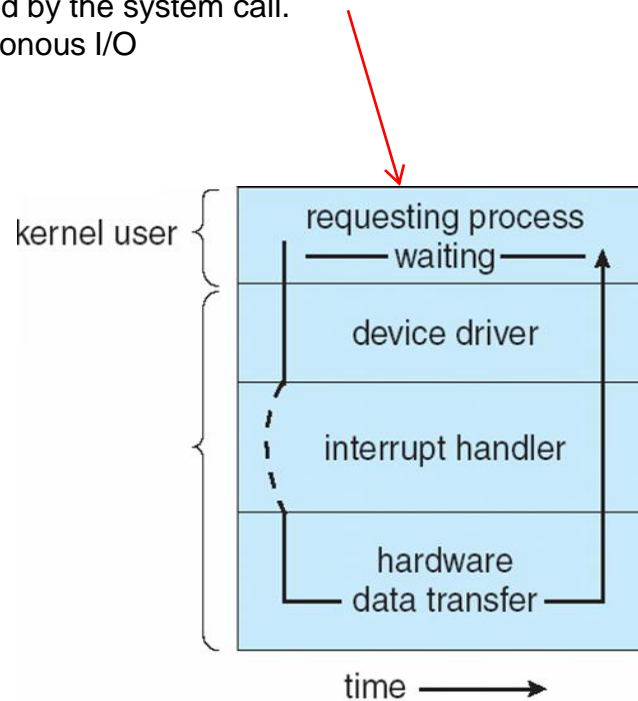
## Clocks and Timers

- Provide current time, elapsed time, timer
- Normal resolution about 1/60 second
- Some systems provide higher-resolution timers
- **Programmable interval timer** used for timings, periodic interrupts
- `ioctl()` (on UNIX) covers odd aspects of I/O such as clocks and timers

# Nonblocking and Asynchronous I/O

## Blocking I/O System Call

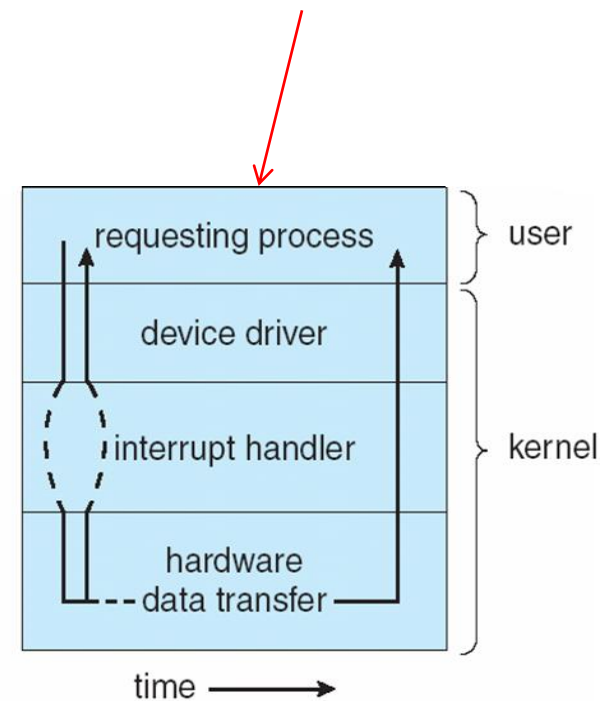
- When an application issues a blocking system call, the execution of the application is suspended. process suspended until I/O completed;
- The application is moved to a wait queue;
- After the system call completes, the application resume.
- When it resume execution, it will receive the values returned by the system call.
- Synchronous I/O



(a)  
Synchronous

## Non-blocking I/O System Call

- Non-blocking I/O system call returns immediately without waiting for the I/O to complete.
- For example, a video application that reads frames from a file on disk while simultaneously decompressing and displaying on the display.
- Asynchronous I/O



(b)  
Asynchronous



# Kernel I/O Subsystem


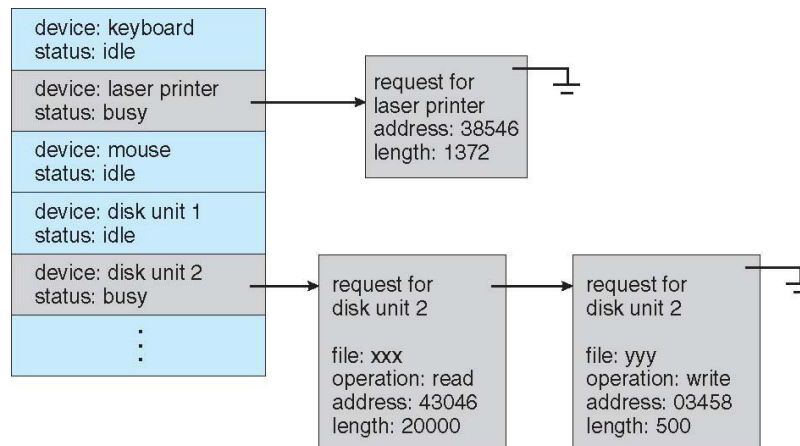
Kernels provide many services related to I/O including

## ■ I/O Scheduling

To determine a good order in which to execute a set of I/O requests

- Operating system implement scheduling by maintaining a wait queue of requests for each device (e.g., disk, CPU)..
- When an application issues a blocking I/O system call, the request is placed on the queue for that device; the I/O scheduler rearranges the order of the queue.
- When a kernel supports asynchronous I/O, it must be able to keep track of many I/O requests at the same time. For this purpose, the operating system might attach the wait queue to a **device – status table**.

- Each I/O device – an entry in table
- Each table entry indicates the device type, address, state; if the device is busy with a request, the type of request and the other parameters will be stored in the table entry.




# Kernel I/O Subsystem (cont.)

## Buffering I/O

- A **buffer** is a memory area that stores data being transferred between two devices or between a device and an application.
- **Buffering is done for three reasons**
  - To cope with device speed mismatch  
E.g., a file is received via modem for storage on the hard disk
  - To cope with device transfer size mismatch  
E.g., a large message is fragmented into small network packets; the packets are sent over the network, and the receiving side places them in a buffer to form the image of the source data.
  - To maintain “copy semantics”  
Application buffer changes while kernel buffer has no effect

Sun Enterprise 6000 Device-Transfer Rates



# Kernel I/O Subsystem (cont.)

## ■ Caching

- A region of fast memory that holds copies of data
- Caching and buffering are distinct functions, but sometimes a region of memory can be used for both purposes.

## ■ Spooling

- A spool is a buffer that holds output for a device, such as a printer, that can serve only one request at a time

## ■ Error Handling

- An operating system uses protected memory to guard against many kinds of hardware and application errors.
- Devices and I/O transfers can fail in many ways
  - Transient reason – network overloaded – can be effectively solved by Operating System
  - Permanent reason – disk controller failure – unlikely to be recovered by Operating System

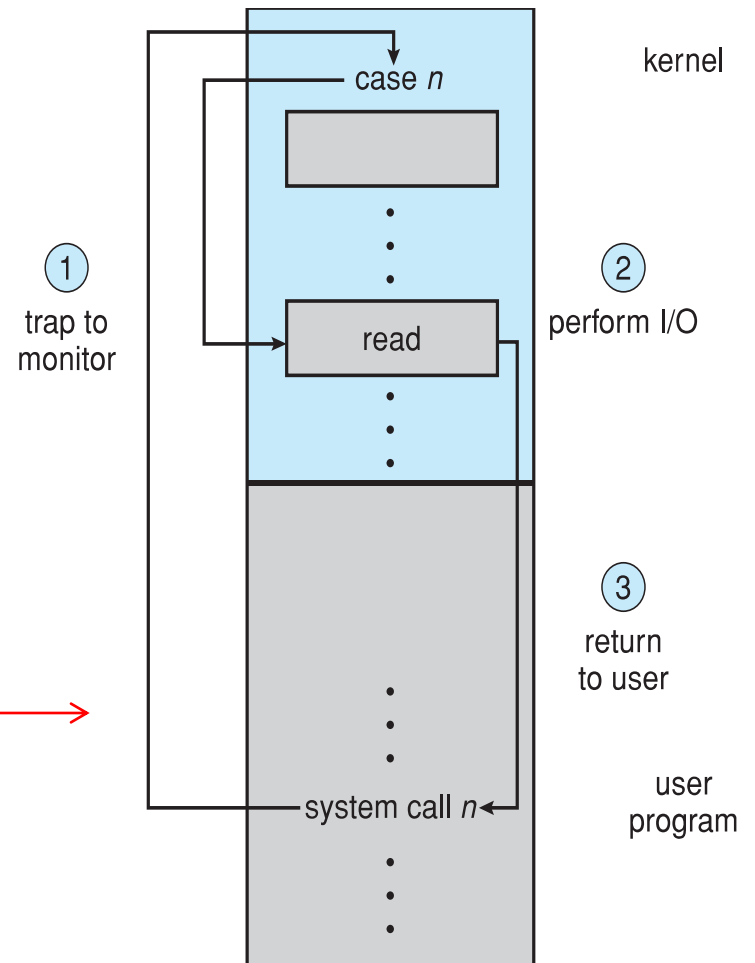
# Kernel I/O Subsystem (cont.)

## ■ I/O Protection

A user process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions. To prevent users from performing illegal I/O

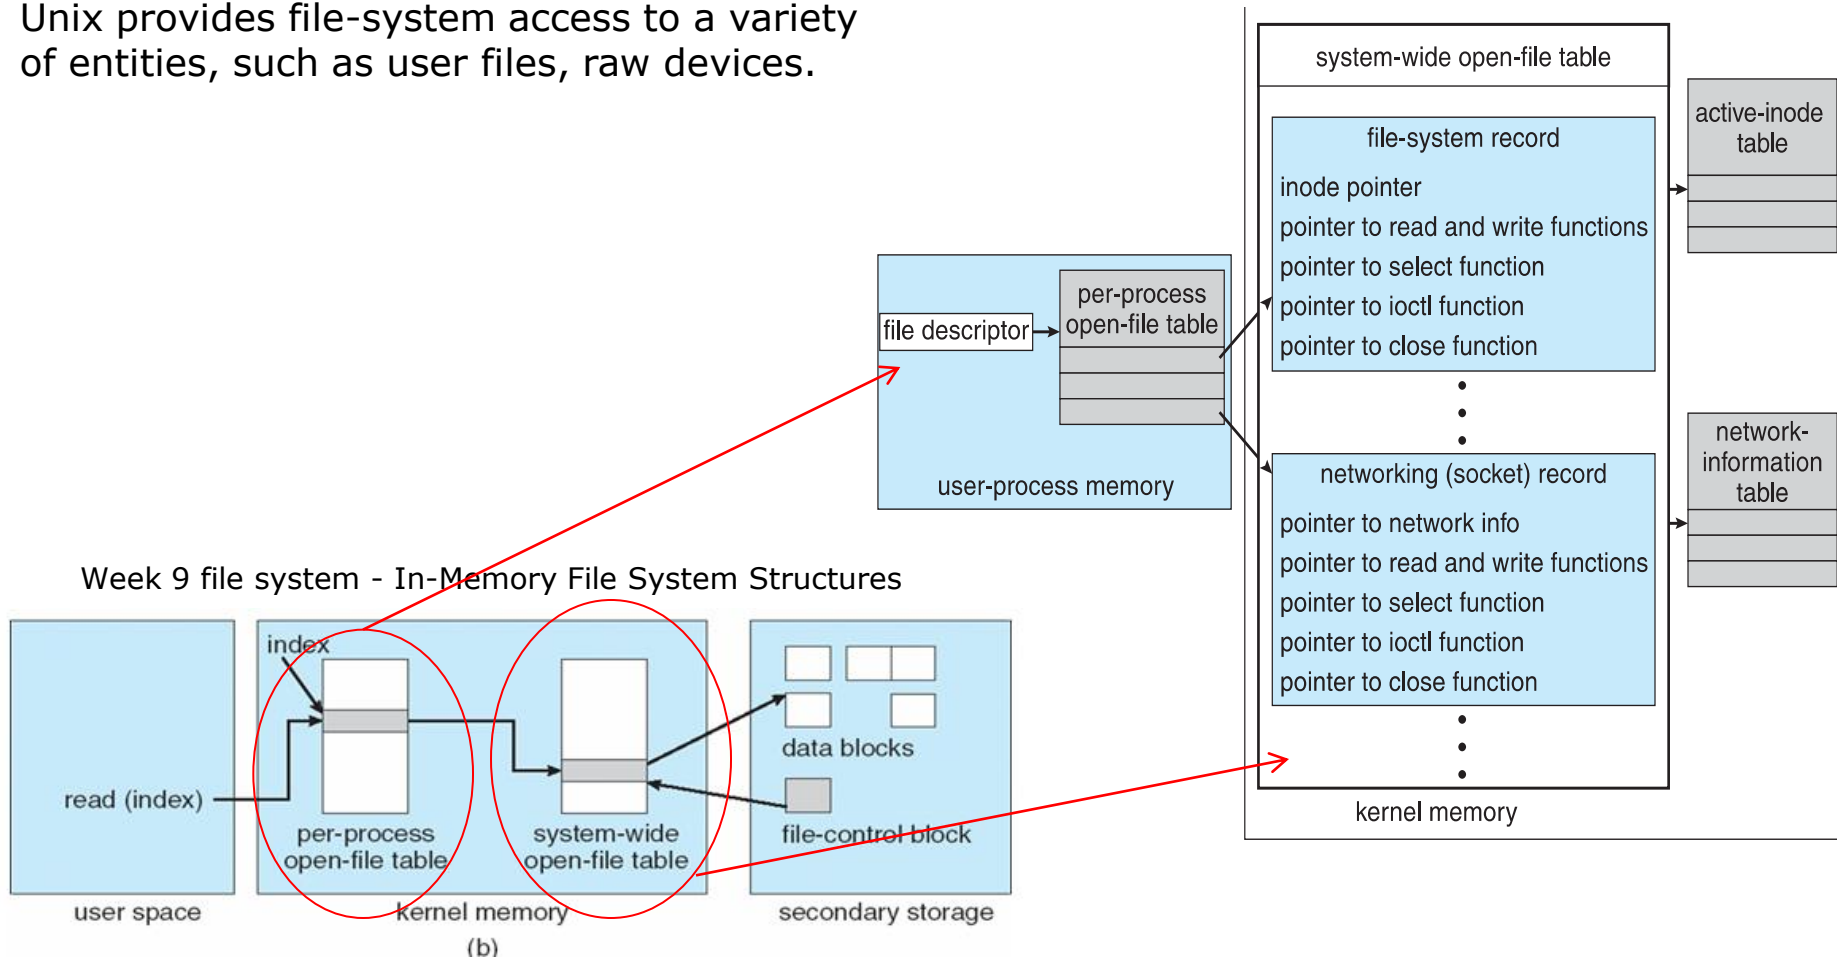
- We define all I/O instructions to be privileged instructions; thus users cannot issue I/O instructions directly; they must do it through the operating system, i.e., **using system call to request the operating system to perform I/O.**
- Memory-mapped and I/O port memory locations must be protected from user access.

## Use of a System Call to Perform I/O



# Kernel Data Structures

- Kernel needs to keep state information of I/O components, including open file tables, network connections, character device state.
  - It does so through a variety of in-kernel data structures, such as the open-file table.
- Unix provides file-system access to a variety of entities, such as user files, raw devices.



# Transforming I/O Requests to Hardware Operations

- Consider reading a file from disk for a process:
  - A process refers to the data by a file name;
  - Within a disk, the file system maps from the file name through the file-system directories to obtain the space allocation of the file;
  - Physically read data from disk into buffer
  - Make data available to the requesting process
  - Return control to process



Life Cycle of An I/O Request

A process issues a blocking read system call to a file that has been opened previously

The system call code checks; if the data is already in the buffer cache, the data are returned to the process and the I/O request is completed.

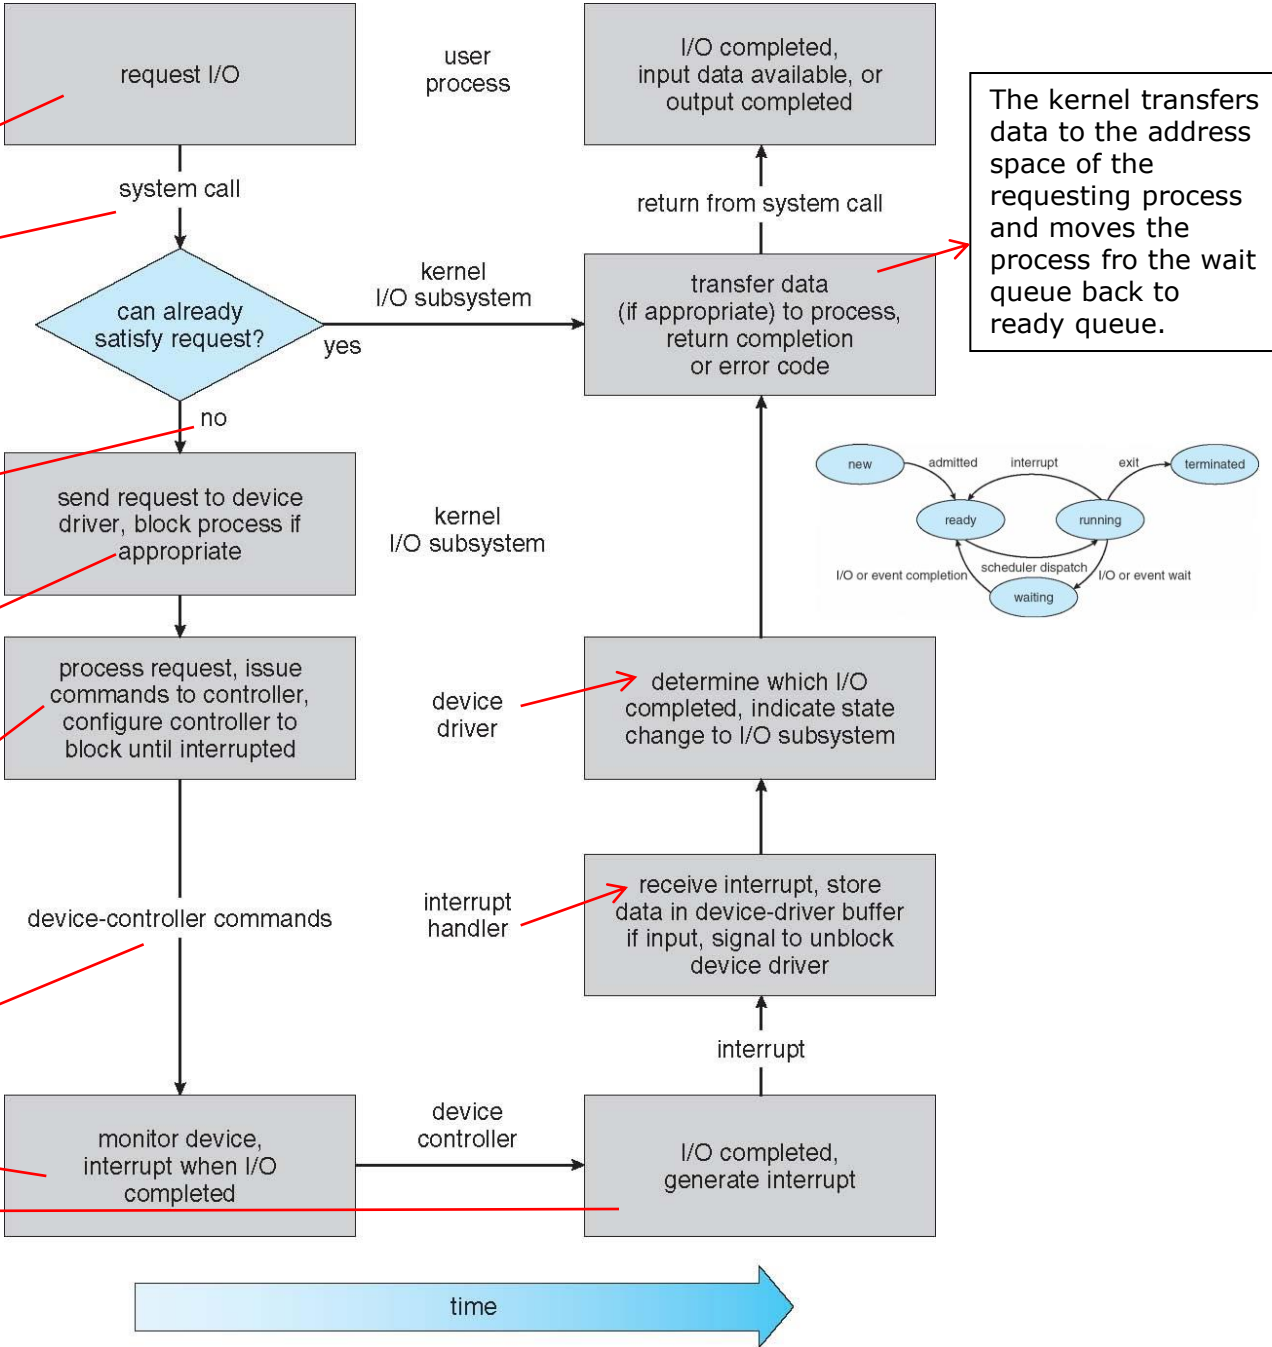
Otherwise, a physical I/O must be performed.

A process is removed to waiting queue

I/O request is scheduled, eventually the I/O subsystem sends the request to the device driver; The device driver allocates kernel buffer space to receive the data.

The device driver sends commands to the device controller by writing into the device-control registers.

Assume DMA manage the data transfer, when the transfer complete, an interrupt is generated.



Next Week

Lecture 11 – Protection

Tutorial 10